

+ Code + Text

```
!pip install -q git+https://github.com/huggingface/transformers
```

```
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
```

```
!curl -L "https://app.roboflow.com/ds/[GqRnfdLze9mzBC06Y4Ps]" > roboflow.zip; unzip roboflow.zip; rm roboflow.zip
```

```
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
100    27    100    27      0      0    128      0  --:--:--  --:--:--  --:--:--  128
```

```
Archive:  roboflow.zip
```

```
End-of-central-directory signature not found.  Either this file is not
a zipfile, or it constitutes one disk of a multi-part archive.  In the
latter case the central directory and zipfile comment will be found on
the last disk(s) of this archive.
```

```
unzip:  cannot find zipfile directory in one of roboflow.zip or
        roboflow.zip.zip, and cannot find roboflow.zip.ZIP, period.
```

```
!pip install roboflow
```

```
from roboflow import Roboflow
rf = Roboflow(api_key="GqRnfdLze9mzBC06Y4Ps")
project = rf.workspace("rmk-engineering-college").project("tumor-vit")
dataset = project.version(1).download("folder")
```

```
Requirement already satisfied: roboflow in /usr/local/lib/python3.10/dist-packages (1.1.0)
Requirement already satisfied: certifi==2022.12.7 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2022.12.7)
Requirement already satisfied: chardet==4.0.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.0.0)
Requirement already satisfied: cyclers==0.10.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (0.10.0)
Requirement already satisfied: idna==2.10 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.10)
Requirement already satisfied: kiwisolver==1.3.1 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.4.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from roboflow) (3.7.1)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.22.4)
Requirement already satisfied: opencv-python>=4.1.2 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.7.0.72)
Requirement already satisfied: Pillow>=7.1.2 in /usr/local/lib/python3.10/dist-packages (from roboflow) (8.4.0)
Requirement already satisfied: pyparsing==2.4.7 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.4.7)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.8.2)
Requirement already satisfied: python-dotenv in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.0.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.27.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.16.0)
Requirement already satisfied: supervision in /usr/local/lib/python3.10/dist-packages (from roboflow) (0.11.1)
Requirement already satisfied: urllib3>=1.26.6 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.26.16)
Requirement already satisfied: wget in /usr/local/lib/python3.10/dist-packages (from roboflow) (3.2)
Requirement already satisfied: tqdm>=4.41.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.65.0)
Requirement already satisfied: PyYAML>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from roboflow) (6.0)
Requirement already satisfied: requests-toolbelt in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.0.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (1.1.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (4.40.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (23.1)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests->roboflow) (2.0.
loading Roboflow workspace...
loading Roboflow project...
Downloading Dataset Version Zip in tumor-ViT-1 to folder: 100% [21405884 / 21405884] bytes
Extracting Dataset Version Zip to tumor-ViT-1 in folder:: 100%[██████████] 2908/2908 [00:00<00:00, 3320.36it/s]
```

```
import torchvision
from torchvision.transforms import ToTensor
```

```
train_ds = torchvision.datasets.ImageFolder('/content/tumor-ViT-1/train', transform=ToTensor())
valid_ds = torchvision.datasets.ImageFolder('/content/tumor-ViT-1/valid', transform=ToTensor())
test_ds = torchvision.datasets.ImageFolder('/content/tumor-ViT-1/test', transform=ToTensor())
```

```
from transformers import ViTModel
from transformers.modeling_outputs import SequenceClassifierOutput
import torch.nn as nn
import torch.nn.functional as F
```

```
class ViTForImageClassification(nn.Module):
    def __init__(self, num_labels=3):
        super(ViTForImageClassification, self).__init__()
        self.vit = ViTModel.from_pretrained('google/vit-base-patch16-224-in21k')
        self.dropout = nn.Dropout(0.1)
        self.classifier = nn.Linear(self.vit.config.hidden_size, num_labels)
        self.num_labels = num_labels

    def forward(self, pixel_values, labels):
        outputs = self.vit(pixel_values=pixel_values)
```

```

output = self.dropout(outputs.last_hidden_state[:,0])
logits = self.classifier(output)

loss = None
if labels is not None:
    loss_fct = nn.CrossEntropyLoss()
    loss = loss_fct(logits.view(-1, self.num_labels), labels.view(-1))
if loss is not None:
    return logits, loss.item()
else:
    return logits, None

EPOCHS = 3
BATCH_SIZE = 10
LEARNING_RATE = 2e-5

from transformers import ViTFeatureExtractor
import torch.nn as nn
import torch
# Define Model
model = ViTForImageClassification(len(train_ds.classes))
# Feature Extractor
feature_extractor = ViTFeatureExtractor.from_pretrained('google/vit-base-patch16-224-in21k')
# Adam Optimizer
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
# Cross Entropy Loss
loss_func = nn.CrossEntropyLoss()
# Use GPU if available
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
if torch.cuda.is_available():
    model.cuda()

Downloading (...)lve/main/config.json: 100% 502/502 [00:00<00:00, 23.3kB/s]
Downloading pytorch_model.bin: 100% 346M/346M [00:03<00:00, 84.8MB/s]
Downloading (...)rocessor_config.json: 100% 160/160 [00:00<00:00, 5.24kB/s]
The class ViTFeatureExtractor is deprecated and will be removed in version 5 of Transformers. Please use ViTImageProcessor instead.

import torch.utils.data as data
from torch.autograd import Variable
import numpy as np

print("Number of train samples: ", len(train_ds))
print("Number of test samples: ", len(test_ds))
print("Detected Classes are: ", train_ds.class_to_idx)

train_loader = data.DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True, num_workers=4)
test_loader = data.DataLoader(test_ds, batch_size=BATCH_SIZE, shuffle=True, num_workers=4)

# Train the model
for epoch in range(EPOCHS):
    for step, (x, y) in enumerate(train_loader):
        # Change input array into list with each batch being one element
        x = np.split(np.squeeze(np.array(x)), BATCH_SIZE)
        # Remove unnecessary dimension
        for index, array in enumerate(x):
            x[index] = np.squeeze(array)
        # Apply feature extractor, stack back into 1 tensor and then convert to tensor
        x = torch.tensor(np.stack(feature_extractor(x)['pixel_values'], axis=0))
        # Send to GPU if available
        x, y = x.to(device), y.to(device)
        b_x = Variable(x) # batch x (image)
        b_y = Variable(y) # batch y (target)
        # Feed through model
        output, loss = model(b_x, None)
        # Calculate loss
        if loss is None:
            loss = loss_func(output, b_y)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    if step % 50 == 0:
        # Get the next batch for testing purposes
        test = next(iter(test_loader))
        test_x = test[0]
        # Reshape and get feature matrices as needed
        test_x = np.split(np.squeeze(np.array(test_x)), BATCH_SIZE)
        for index, array in enumerate(test_x):

```

```

    test_x[index] = np.squeeze(array)
test_x = torch.tensor(np.stack(feature_extractor(test_x)['pixel_values'], axis=0))
# Send to appropriate computing device
test_x = test_x.to(device)
test_y = test[1].to(device)
# Get output (+ respective class) and compare to target
test_output, loss = model(test_x, test_y)
test_output = test_output.argmax(1)
# Calculate Accuracy
accuracy = (test_output == test_y).sum().item() / BATCH_SIZE
print('Epoch: ', epoch, '| train loss: %.4f' % loss, '| test accuracy: %.2f' % accuracy)

```

```

Number of train samples: 2300
Number of test samples: 350
Detected Classes are: {'glioma': 0, 'meningioma': 1, 'notumor': 2, 'pituitary': 3}
This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller
Epoch: 0 | train loss: 1.4159 | test accuracy: 0.10
Epoch: 0 | train loss: 1.3497 | test accuracy: 0.20
Epoch: 0 | train loss: 1.3774 | test accuracy: 0.40
Epoch: 0 | train loss: 1.3818 | test accuracy: 0.30
Epoch: 0 | train loss: 1.3927 | test accuracy: 0.00
Epoch: 1 | train loss: 1.3881 | test accuracy: 0.30
Epoch: 1 | train loss: 1.4104 | test accuracy: 0.00
Epoch: 1 | train loss: 1.4370 | test accuracy: 0.00
Epoch: 1 | train loss: 1.3937 | test accuracy: 0.30
Epoch: 1 | train loss: 1.3987 | test accuracy: 0.20
Epoch: 2 | train loss: 1.3948 | test accuracy: 0.10
Epoch: 2 | train loss: 1.3916 | test accuracy: 0.20
Epoch: 2 | train loss: 1.4005 | test accuracy: 0.10
Epoch: 2 | train loss: 1.3808 | test accuracy: 0.30
Epoch: 2 | train loss: 1.3918 | test accuracy: 0.20

```

```

import matplotlib.pyplot as plt
import numpy as np

```

```

EVAL_BATCH = 1
eval_loader = data.DataLoader(valid_ds, batch_size=EVAL_BATCH, shuffle=True, num_workers=4)
# Disable grad
with torch.no_grad():

    inputs, target = next(iter(eval_loader))
    # Reshape and get feature matrices as needed
    print(inputs.shape)
    inputs = inputs[0].permute(1, 2, 0)
    # Save original Input
    originalInput = inputs
    for index, array in enumerate(inputs):
        inputs[index] = np.squeeze(array)
    inputs = torch.tensor(np.stack(feature_extractor(inputs)['pixel_values'], axis=0))

    # Send to appropriate computing device
    inputs = inputs.to(device)
    target = target.to(device)

    # Generate prediction
    prediction, loss = model(inputs, target)

    # Predicted class value using argmax
    predicted_class = np.argmax(prediction.cpu())
    value_predicted = list(valid_ds.class_to_idx.keys())[list(valid_ds.class_to_idx.values()).index(predicted_class)]
    value_target = list(valid_ds.class_to_idx.keys())[list(valid_ds.class_to_idx.values()).index(target)]

    # Show result
    plt.imshow(originalInput)
    plt.xlim(224,0)
    plt.ylim(224,0)
    plt.title(f'Prediction: {value_predicted} - Actual target: {value_target}')
    plt.show()

```

```

torch.Size([1, 3, 224, 224])
Prediction: pituitary - Actual target: pituitary
0
25
50
75
100

torch.save(model, '/content/ViTmodel1.pt')

torch.save(model, '/home/ViTmodelsaved.pt')

MODEL_PATH = '/content/ViTmodel1.pt'
model = torch.load(MODEL_PATH)
model.eval()

ViTForImageClassification(
  (vit): ViTModel(
    (embeddings): ViTEmbeddings(
      (patch_embeddings): ViTPatchEmbeddings(
        (projection): Conv2d(3, 768, kernel_size=(16, 16), stride=(16, 16))
      )
      (dropout): Dropout(p=0.0, inplace=False)
    )
    (encoder): ViTEncoder(
      (layer): ModuleList(
        (0-11): 12 x ViTLayer(
          (attention): ViTAttention(
            (attention): ViTSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.0, inplace=False)
            )
            (output): ViTSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.0, inplace=False)
            )
          )
          (intermediate): ViTIntermediate(
            (dense): Linear(in_features=768, out_features=3072, bias=True)
            (intermediate_act_fn): GELUActivation()
          )
          (output): ViTOutput(
            (dense): Linear(in_features=3072, out_features=768, bias=True)
            (dropout): Dropout(p=0.0, inplace=False)
          )
          (layernorm_before): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (layernorm_after): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        )
      )
    )
    (layernorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (pooler): ViTPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.1, inplace=False)
  (classifier): Linear(in_features=768, out_features=4, bias=True)
)

from PIL import Image
image_path = '/Hirnmetastase_MRT-T1_KM.jpg'
image = Image.open(image_path)

import torch
from torchvision import transforms
from PIL import Image
import matplotlib.pyplot as plt

# Load and preprocess the image

preprocess = transforms.Compose([
  transforms.Resize((224, 224)),
  transforms.ToTensor(),
  transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

```

```
input_image = preprocess(image).unsqueeze(0)

# Send the input image to the appropriate device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
input_image = input_image.to(device)

# Disable grad
with torch.no_grad():
    # Generate prediction
    logits, _ = model(input_image, labels=None)

# Convert prediction to class label
predicted_class_index = torch.argmax(logits, dim=1).item()
predicted_class_label = list(valid_ds.class_to_idx.keys())[predicted_class_index]

# Show result
print(predicted_class_label)
plt.imshow(image)
plt.title(f'Prediction: {predicted_class_label}')
plt.show()
```

